

6.1220 LECTURE 21

BRENDON JIANG

CONTENTS

1	Sketching and Streaming	1
1.1	Counting	1
1.2	Distinct Elements	3
1.3	Similarity Search	5

1. SKETCHING AND STREAMING

1.1. Counting

Suppose we are given a dataset X , we want to compress X into a sketch $S(X)$ such that $f(X)$ can be computed/approximated given only $S(X)$. X will be given via a data stream which you get to read one-step at a time. We will need to update $S(X)$ "on-the-fly".

Let us start with a simple example.

Example 1.1

Data Stream: $X = 1, 1, 1, \dots$ of some unknown length n . The goal is to compute $f(X) = n = \# \text{ of } 1\text{s}$.

Solution 1.2

The easy algorithm is to just maintain a counter $S = \# \text{ of } 1\text{s}$ seen so far. The space used is then $\log n$. Note that this algorithm computes the optimal solution.

What if we want $o(\log n)$ space and an approximation instead?

Solution 1.3: First Attempt

Increment the counter S every k th 1 for constant k . At the end, output $k \cdot S$. However, we didn't really save space here, as the space is $\log(n) - \log(k)$.

Example 1.4: Second Attempt

Instead of counting n , we approximate n by counting $\log n$. However, currently it is unclear when we should increment $S(X)$. We can get around this through randomization, and achieve $O(\log \log n)$ space until return.

Algorithm 1: Morris' Algorithm

```

1 Initialize  $S \leftarrow 0$ 
2 while stream is not empty do
3   Sample  $b \in \{0, 1\}$  where  $b = 1$  with probability  $\frac{1}{2^S}$ .
4   if  $b=1$  then
5     Increment  $S \leftarrow S + 1$ 
6 return  $2^S - 1$ .
```

Now let us prove this algorithm is actually a good approximation. Suppose $\tilde{f}(n) = 2^S - 1$ for S produced from running this algorithm on a length- n stream of 1s.

Lemma 1.5

$\forall n \in \mathbb{N}, \mathbb{E}[\tilde{f}(n)] = n$ and $\text{Var}(\tilde{f}(n)) < \frac{1}{2}n^2$.

Proof

We will prove by induction. In the base case, we note that $\tilde{f}(0) = 0, \tilde{f}(1) = 1$ with probability 1. In the inductive step, we find

$$\begin{aligned}
\mathbb{E}[2^{S(k+1)}] &= \sum_{j=0}^k \mathbb{E}[2^{S(k+1)} | S(k) = j] \cdot \Pr[S(k) = j] \\
&= \sum_{j=0}^k \left(\frac{1}{2^j} \cdot 2^{j+1} + \left(1 - \frac{1}{2^j}\right) \cdot 2^j \right) \cdot \Pr[S(k) = j] \\
&= \sum_{j=0}^k (2^j + 1) \cdot \Pr[S(k) = j] \\
&= \sum_{j=0}^k 2^j \cdot \Pr[S(k) = j] + \sum_{j=0}^k \Pr[S(k) = j] \\
&= k + 1 + 1
\end{aligned}$$

A similar proof establishes the variant part of the lemma.

Corollary 1.6

$$\Pr[|\tilde{f}(n) - n| > \epsilon] \leq \frac{1}{2\epsilon^2}$$

Proof

Use Chebyshev.

Corollary 1.7

$$\forall n \in \mathbb{N}, \Pr[(1 - \epsilon)n \leq \tilde{f}(n) \leq (1 + \epsilon)n] \geq 1 - \frac{1}{2\epsilon^2}.$$

To boost the accuracy and success probability of our algorithm, we can simply run it a bunch of times in the same data stream and return the median of the mean.

1.2. Distinct Elements

Data Stream: Length- n sequence X of elements from some universe \mathcal{U} . Our goal is to estimate $d = d(X) = \#$ of distinct elements occurring in the stream X . Note that there are two easy algorithms, which is to maintain a boolean array and a container of the distinct elements respectively.

Algorithm 2: Flajolet-Martin Algorithm

```

7 Initialize  $h : \mathcal{U} \rightarrow [0, 1]$  where  $h(u) \sim \text{Unif}[0, 1]$  independently for all  $u \in \mathcal{U}$  Initialize  $S \leftarrow 1$ 
8 while stream is not empty: do
9   | Let  $x$  be the next stream element Update  $S \leftarrow \min\{S, h(x)\}$ 
10 return  $\frac{1}{S} - 1$ 

```

Lemma 1.8

Let $S(X)$ denote the sketch produced given a sequence X .

$$\forall X \in \mathcal{U}^*, \mathbb{E}[S(x)] = \frac{1}{d+1} \text{ and } \text{Var}(S(x)) = \frac{d}{(d+1)^2}(d+2)$$

Proof

Using layered cake representation, we find

$$\begin{aligned}
LHS &= \int_0^1 \Pr[\min\{u_1, \dots, u_d\} \geq y] dy \\
&= \int_0^1 \Pr[u_i \geq y \forall i] dy \\
&= \int_0^1 \Pr[u_i \geq y]^d dy \\
&= \int_0^1 (1 - y)^d dy \\
&= \frac{1}{d+1}
\end{aligned}$$

Again, we can boost the probability of this algorithm by running it multiple times and merging the answers.

Right now, there are two problems with this algorithms. We can efficiently store a truly uniformly random hash function $h : \mathcal{U} \rightarrow [0, 1]$. We also want h to map to a discrete set instead. Recall that

Definition 1.9

A family of hash functions is 2-wise independent if $\forall x, y \in \mathcal{U}$, and $\forall s, t \in \{0, \dots, m-1\}$,

$$\mathbb{P}_{h \sim \mathcal{H}}[h(x) = s, h(y) = t] = \frac{1}{m^2}$$

Algorithm 3: k-Minimum Value (KMV) Algorithm:

```

11 Initialize  $h : \mathcal{U} \rightarrow \{0, \dots, m-1\}$  from 2-wise independent hash family  $\mathcal{H}$ 
12 Initialize  $S =$ 
13 Initialize  $k = \frac{24}{\epsilon^2}$ 
14 while stream is not empty: do
15   | Let  $x$  be the next stream element
16   | Update  $S \leftarrow \{k \text{ smallest elements of } S \cup \{h(x)\}\}$ 
17 if  $|S| = k$  then
18   | return  $\frac{km}{\max(S)}$ 
19 else
20   | return  $|S|$ 

```

This algorithm has space $O(\frac{1}{\epsilon^2} \log |\mathcal{U}|)$. Update Time: $O(\log(|\mathcal{U}|) \cdot \log(\frac{1}{\epsilon}))$

1.3. Similarity Search

We are now not using the streaming model. The input to this problem is a collection of sets $A_1, \dots, A_n \subseteq \mathcal{U}$ where $|A_i| \leq d$ for all $i \in [n]$. The queries we want is for a new set $A \subseteq \mathcal{U}$ satisfying $|A| \leq d$ and $s \in [0, 1]$, does there exist A_i such that the Jaccard similarity satisfies $J(A, A_i) \geq s$? Here, we define

$$J(A, A_i) := \frac{|A \cap A_i|}{|A \cup A_i|}$$

The easy algorithm is to just compute $J(A, A_i)$ exactly for each $i \in n$, and obtain a time bound of $O(nd)$.

Our goal is to construct a hash function $\sigma : 2^{\mathcal{U}} \rightarrow [0, 1]$ such that similar sets collide. Then you just search within the bucket indexed by $\sigma(A)$. We can perhaps build a perfect static hash table with key-value pairs: Key = $\sigma(A_i)$, Value = A_i .

Lemma 1.10: Min-Hash Lemma

For a uniformly random $h : \mathcal{U} \rightarrow [0, 1]$, define the min-hash

$$\sigma_h(A) := \min_{x \in A} h(x)$$

Then for all sets $A, B \subseteq \mathcal{U}$,

$$\Pr_h[\sigma_h(A) = \sigma_h(B)] = J(A, B)$$

Proof

The function h induces an ordering π on \mathcal{U} where $x < y$ iff $h(x) < h(y)$. This means

$$\begin{aligned} LHS &= \Pr[\min \text{ of } A \text{ under } \pi = \min \text{ of } B \text{ under } \pi] \\ &= \Pr_{\pi}[\{\min \text{ of } A \cup B \text{ with respect to } \pi\} \in A \cap B] \\ &= \frac{|A \cap B|}{|A \cup B|} \end{aligned}$$

We want to boost the probability. If $J(A, B) \geq s$, boost the probability that $\sigma(A) = \sigma(B)$. If $J(A, B) < \eta < s$, we reduce the probability that $\sigma(A) = \sigma(B)$ (so that we don't have to search through junk).